



Zero to Viral: Build TikTok-style Recommendation

[linkedin.com/in/yodiaditya](https://www.linkedin.com/in/yodiaditya)



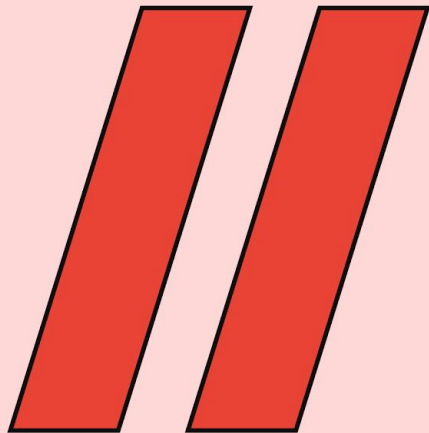
Disclaimer: These are the personal opinions of the author(s). Any assumptions, opinions stated here are theirs and not representative of or attributable to their current or any prior employer(s). Apart from publicly available information, any other information here is not claimed to refer to any company including ones the author(s) may have worked in or been associated with.



Disclaimer

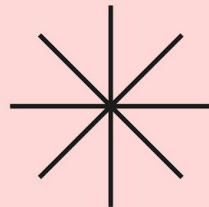
These are the personal opinions of the author(s). Any assumptions, opinions stated here are theirs and not representative of or attributable to their current or any prior employer(s). Apart from publicly available information, any other information here is not claimed to refer to any company including ones the author(s) may have worked in or been associated with.

AI
@DevFest



Google
Developer
Groups

**How do you build short
videos recommendation
when all have you just raw
files with no metadata and
nothing else ?**



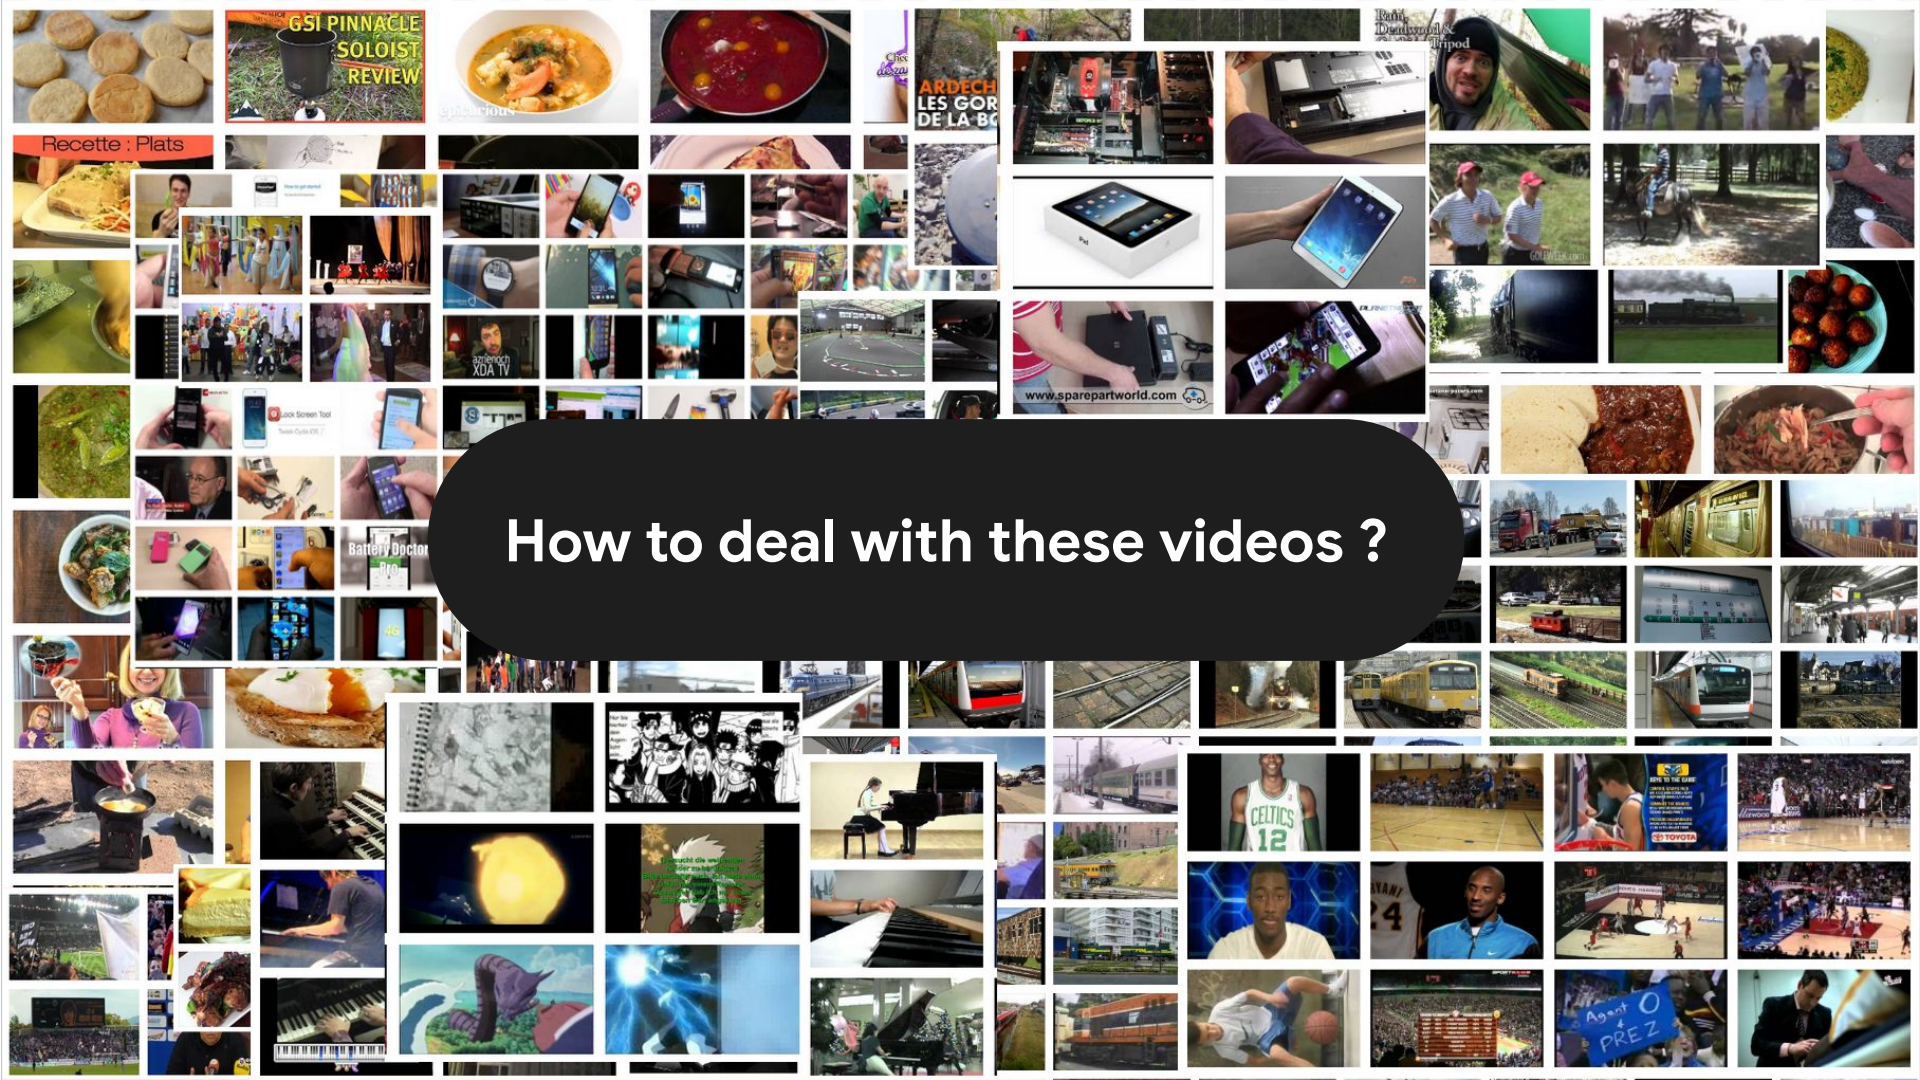
Statistics

Countries with the Most TikTok Users in the World 2024



Estimated 1 Billion
views per day

34 million TikTok
videos are
uploaded daily



How to deal with these videos ?

What you will get in next 17 minutes

(1)

Intro: Recommendation System

(2)

Architecture: Raw Dataset to Final Results

(3)

Code (Yes, its code!)

(4)

Tips to Scale

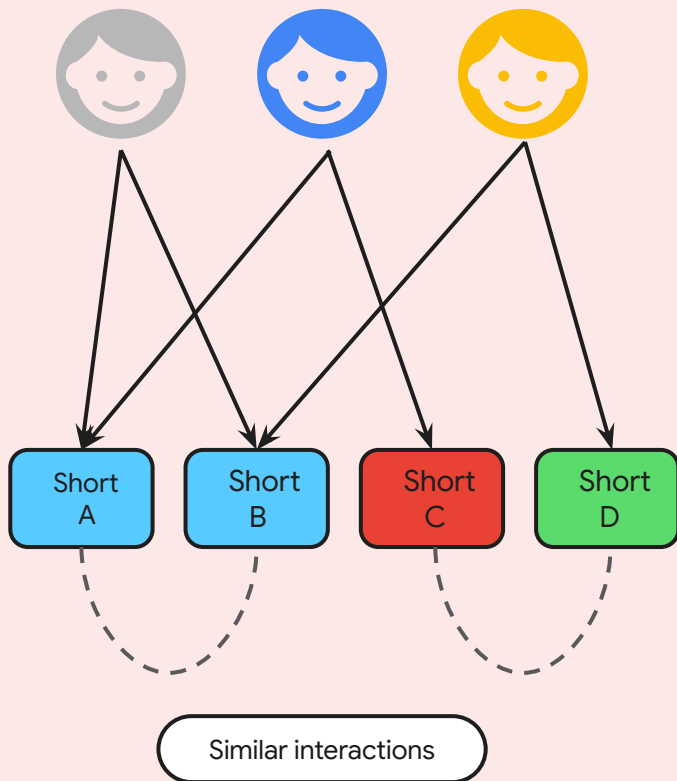
A bit about myself

- Build short-videos recommendation at scale in production serving millions of users
- Certified Google Cloud Professional Machine Learning Engineer (PMLE).
- 15 years working experience multi-domains and loves open-source !

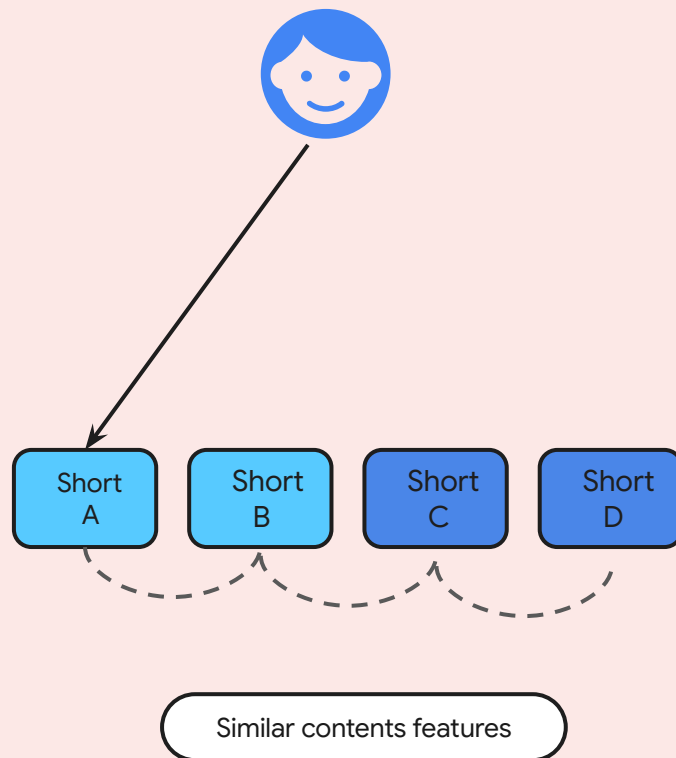


<https://github.com/yodiaditya>

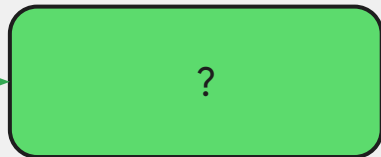
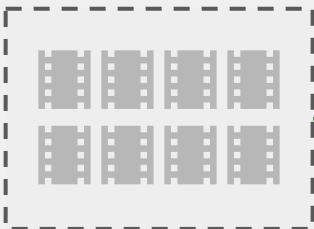
Collaborative Filtering



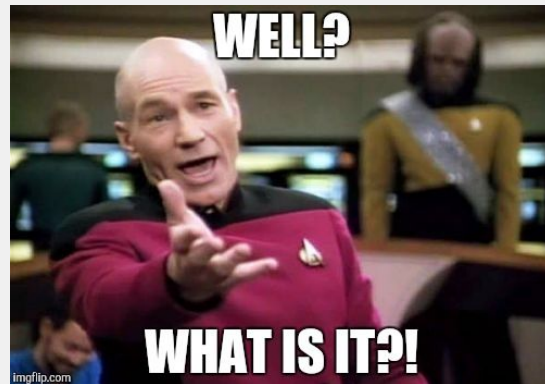
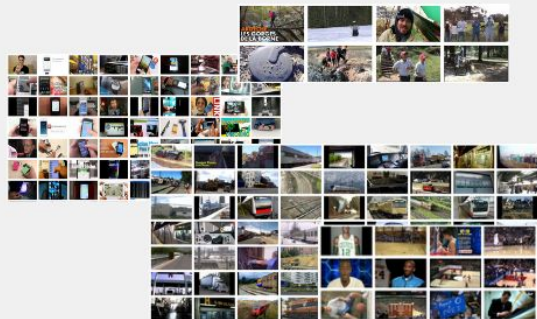
Content Based Filtering



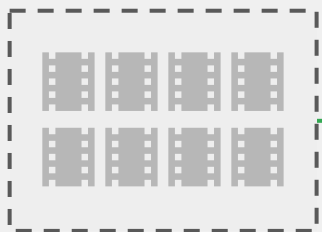
Users upload videos to server



Shorts Videos
(Raw version)



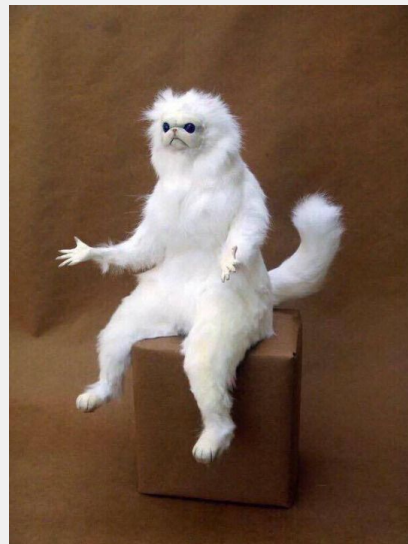
First step is to extract the information from the videos



**Feature
Extraction**

Images

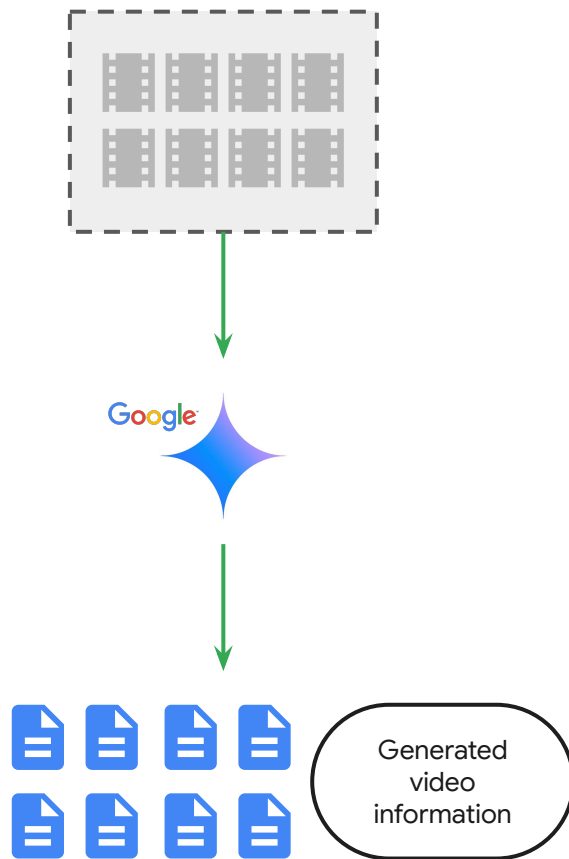
Audio



Gemini: Video Understanding

Model	Video modality details
Gemini 1.5 Flash Go to the Gemini 1.5 Flash model card	<p>Maximum video length:</p> <ul style="list-style-type: none">• With audio: ~50 minutes• Without audio: 60 minutes <p>Maximum videos per prompt: 10</p>
Gemini 1.5 Pro Go to the Gemini 1.5 Pro model card	<p>Maximum video length:</p> <ul style="list-style-type: none">• With audio: ~50 minutes• Without audio: 60 minutes <p>Maximum videos per prompt: 10</p>
Gemini 1.0 Pro Vision Go to the Gemini 1.0 Pro Vision model card	<p>Maximum video length: 2 minutes</p> <p>The maximum videos per prompt: 1</p> <p>Audio in the video is ignored.</p>

<https://cloud.google.com/vertex-ai/generative-ai/docs/multimodal/video-understanding>



Next, Embedding Space

Supported models

You can get text embeddings by using the following models:

English models	Multilingual models
<code>textembedding-gecko@001</code>	<code>textembedding-gecko-multilingual@001</code>
<code>textembedding-gecko@003</code>	<code>text-multilingual-embedding-002</code>
<code>text-embedding-004</code>	
<code>text-embedding-005</code>	

If you are new to these models, we recommend that you use the latest versions. For English text, use `text-embedding-005`. For multilingual text, use `text-multilingual-embedding-002`.

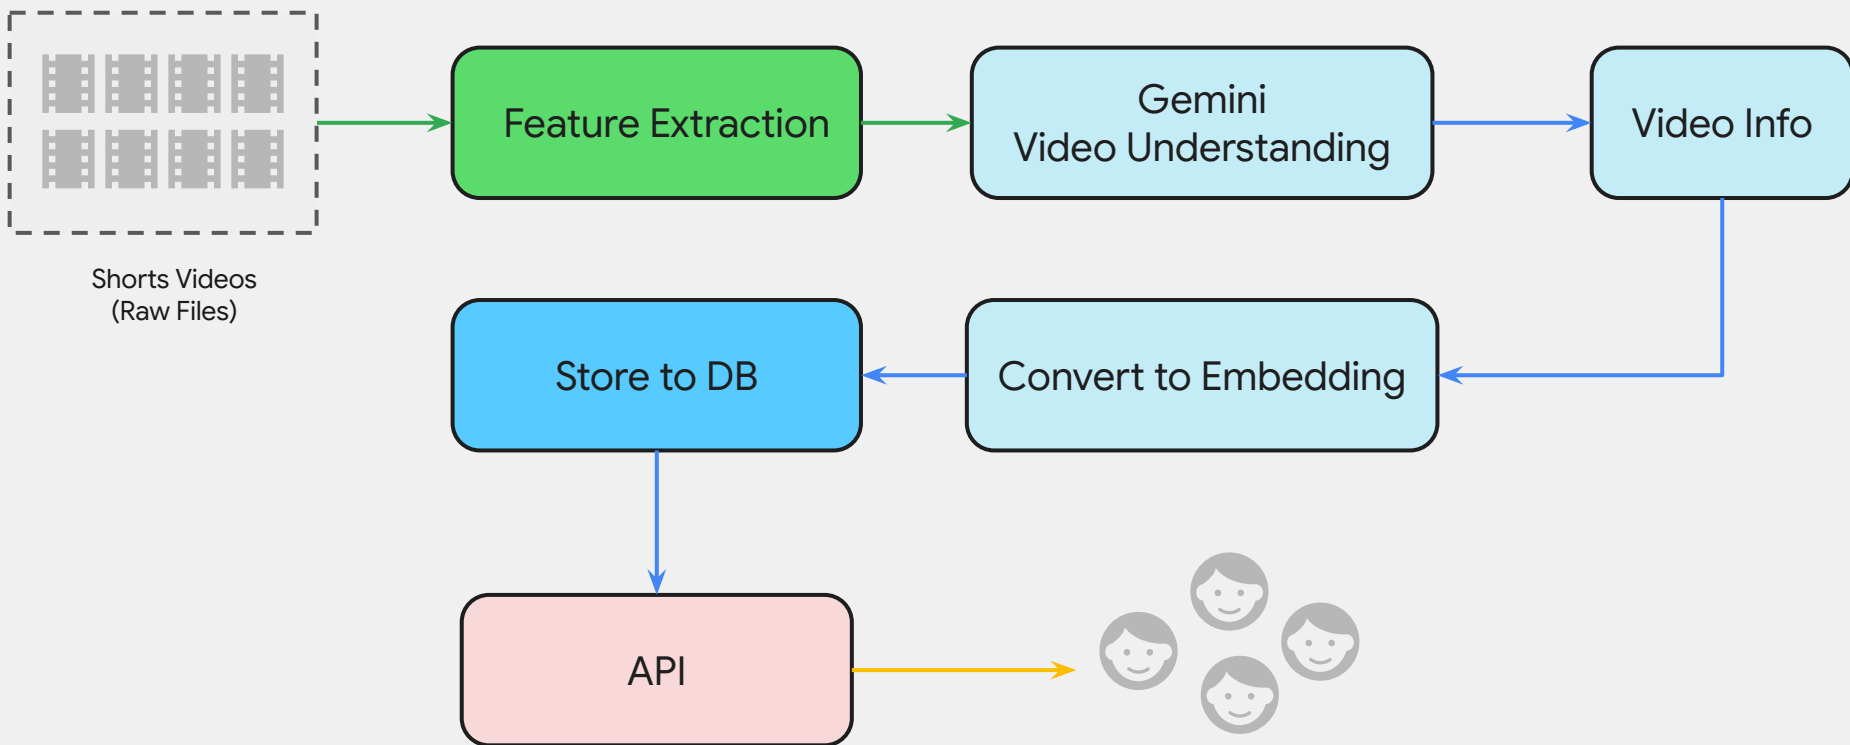
<https://cloud.google.com/vertex-ai/generative-ai/docs/embeddings/get-text-embeddings>



Generated
video
information



Architecture Content-Based Raw Shorts Videos

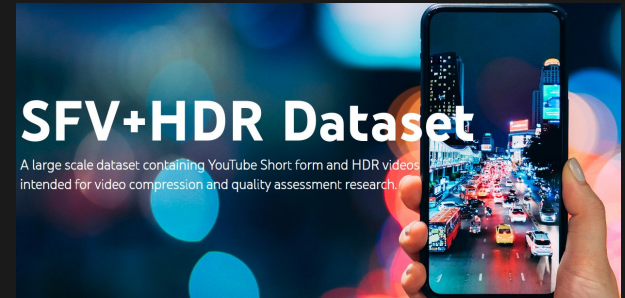


Show me the code

```
import vertexai
from vertexai.generative_models import GenerativeModel, Part

PROJECT_ID = 'YOUR-PROJECT-ID'
vertexai.init(project=PROJECT_ID, location="us-central1")
vision_model = GenerativeModel("gemini-1.5-flash-002")

# Generate text
response = vision_model.generate_content(
    [
        Part.from_uri(
            "gs://shorts-hdr-dataset/videos/sdr/SDR_Animal_45j4.mp4", mime_type="video/mp4"
        ),
        "Watch each frame in video. Do not make up any information that is not part of the video. Generate title of video, genres, short description and taggings of content ",
    ]
)
print(response.text)
```



<https://media.withyoutube.com/sfv-hdr>

Here's a breakdown of the video based on your request:

****Title:**** Fluffy Persian Cat

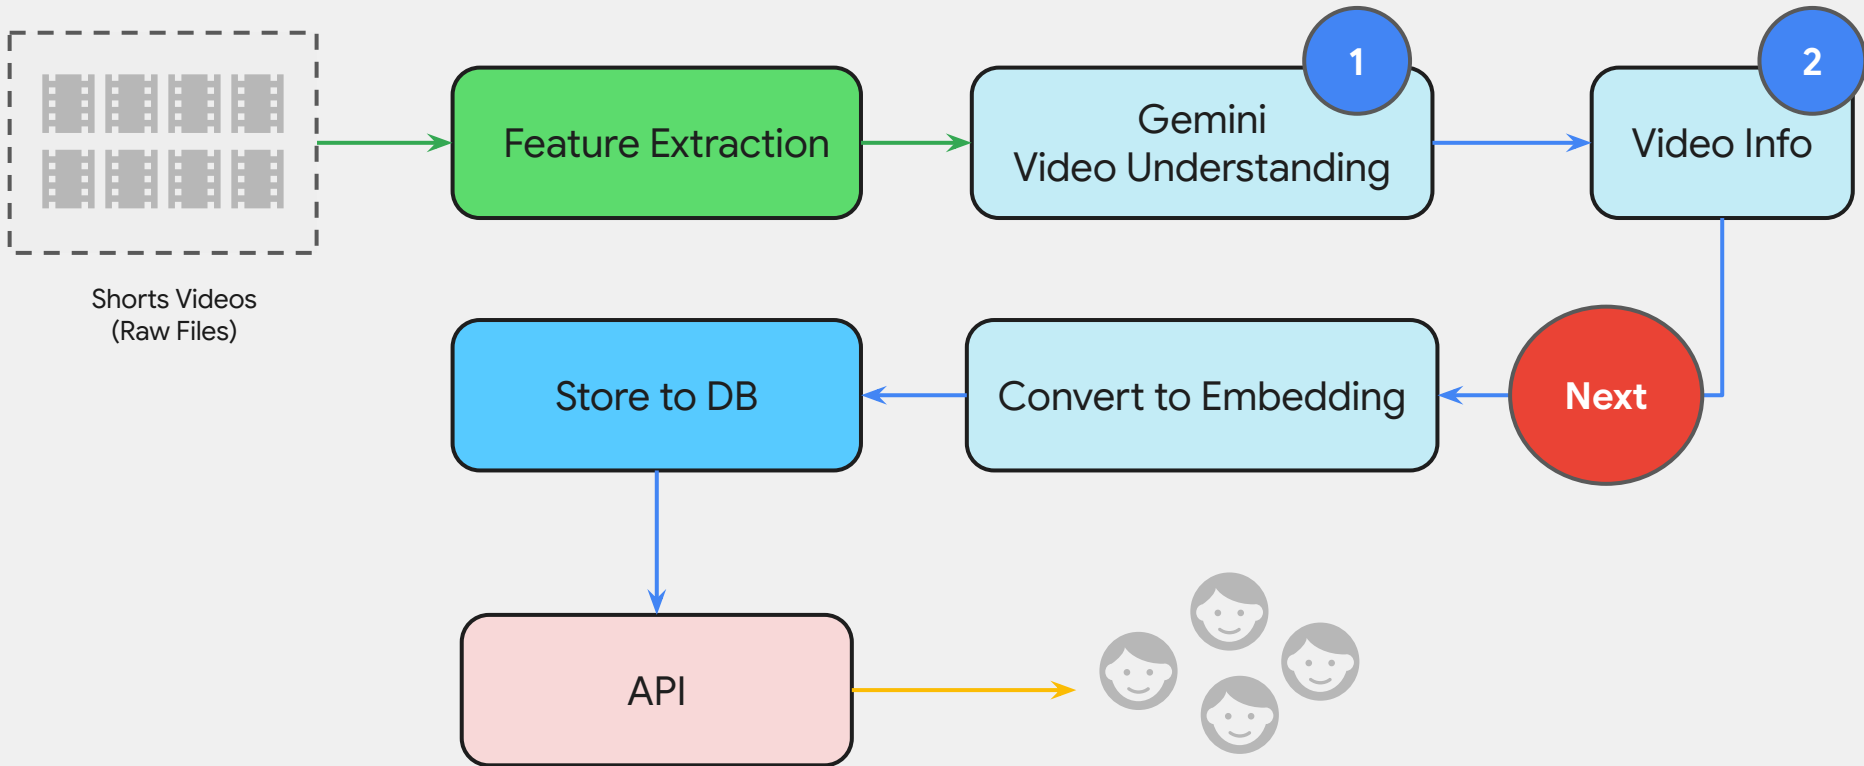
****Genres:**** Animals, Cats, Pets

****Short Description:**** A beautiful fluffy Persian cat sits on a wooden pallet.

****Taggings:**** Persian cat, fluffy cat, cat, kitten, pets, animals, cute, adorable, white cat, long hair cat



Architecture Content-Based Raw Shorts Videos



```
import google.generativeai as genai
import os
```

```
genai.configure(api_key=os.environ["GEMINI_API_KEY"])
```

```
text = f"""
    *Title:** Fluffy Persian Cat **Genres:** Animals, Cats, Pets
    **Short Description:** A beautiful fluffy Persian cat sits on a wooden pallet.
    **Taggings:** Persian cat, fluffy cat, cat, kitten, pets, animals, cute, adorable, white cat, long
    hair cat"""
```

```
result = genai.embed_content(
    model="models/text-embedding-004",
    content=text)
```

```
print(str(result['embedding']))
```


[-0.07675823, 0.027660733, 0.049145754, -0.045824636, 0.061620172, 0.032313757, -0.015517017, -0.03240082, 0.015460015, 0.062071748, -0.03166332, 0.036584534, 0.024837952, 0.009332091, 0.04251558, -0.029856514, 0.05006616, 0.073106274, -0.029902106, -0.025679907, 0.084334895, -0.017889986, 0.059542187, 0.0013102485, -0.046999988, -0.007112105, 0.055407282, 0.018771457, 0.07350461, -0.01813061, 0.04455579, 0.064248785, 0.00448263, 0.008756887, 0.047846865, -0.025803361, -0.019773666, 0.041618045, 0.051186264, -0.07005315, -0.033514354, -0.026876828, -0.0008525868, 0.08978644, 0.027724743, 0.010746625, -0.061409652, 0.026029717, -0.0034387638, 0.035402797, 0.007589881, -0.023214854, -0.019220363, 0.02619021, -0.026539033, 0.012508128, 0.0040561967, -0.028133305, -0.016629174, 0.020951482, 0.06473735, 0.034282178, 0.051744517, 0.019634075, -0.016341783, 0.021416675, -0.05009744, -0.0076771406, -0.036871992, 0.095764354, -0.02469845, 0.025160229, -0.046339314, 0.024582243, -0.008852839, -0.010245625, 0.011931524, 0.037149847, -0.027476897, 0.021095295, -0.0048002778, 0.012314561, 0.09466355, -0.028999066, -0.010545491, 0.019936696, -0.0075119766, 0.02987802, -0.08925607, 0.0030204963, 0.06398722, -0.005665603, 0.055336952, 0.031011213, 0.0803204, -0.046087332, -0.09462911, -0.029559113, 0.017180113, 0.035609562, -0.041688204, 0.016256401, -0.018201571, -0.06340068, 0.05076339, 0.00789683, -0.04663778, 0.003865695, -0.02993782, 0.05526901, -0.059481055, -0.08222241, -0.044425923, 0.03254365, 0.052314047, 0.077546, 0.0024010574, 0.030455906, -0.017964754, -0.03830613, 0.042473085, 0.021389754, 0.031468235, -0.0121501945, -0.010297442, 0.012780487, -0.0086625, -0.019988453, 0.017340286, -0.009815917, 0.068223, 0.02263392, 0.02600014, 0.03895912, -0.01126567, 0.03847954, 0.035594787, -0.08233705, 0.007975588, -0.010537024, -0.042857118, -0.037623975, -0.04778282, 0.007993469, 0.0021377627, -0.046956647, -0.019428788, 0.068178736, -0.002273978, 0.032764167, 0.022261541, -0.050062243, -0.0049963817, -0.0493768, 0.027171357, 0.030386483, 0.035637327, -0.026583953, -0.05276558, -0.051282965, 0.060191035, -0.03913675, 0.009324002, -2.8839599e-05, -0.010929866, -0.0029119896, 0.013345155, -0.013394253, 0.006282062, -0.029159687, -0.023629908, -0.04224124, -0.045065757, -0.023713488, -0.038271822, 0.046776716, -0.024885133, -0.05454553, -0.07411608, 0.02653326, 0.01615619, -0.017124163, -0.03665412, -0.07935555, 0.034953605, -0.073878795, -0.035953894, -0.019765098, 0.004043873, 0.024911957, -0.005691729, -0.050333932, -0.000985027, -0.036075074, 0.011225469, -0.016764432, 0.09254729, -0.072244085, -0.037920788, 0.0124474, 0.041128416, -0.018892227, 0.02873513, 0.03394367, 0.005887189, 0.04828047, -0.015023347, -0.010894143, 0.003211685, -0.042894274, -0.110083245, 0.013853838, 0.02313884, -0.072523855, 0.024407636, 0.02207754, 0.027719164, -0.01624695, 0.040581483, -0.039446704, 0.062565975, 0.014453987, 0.021910533, -0.019675337, 0.04837492, -0.011042686, -0.02570678, -0.0035492862, -0.012308466, -0.020631982, -0.035574436, 0.042452924, 0.022002371, -0.012221022, -0.039625946, 0.009512582, 0.0071309633, -0.01212419, -0.014330538, 0.06173416, -0.007517297, 0.016374068, 0.08243889, 0.0066845445, 0.053010326, 0.004689015, 0.027499422, 0.05117024, -0.06105194, -0.0407715, -0.006838235, 0.007643424, -0.028924558, 0.031687874, -0.06336678, -0.01165728, -0.003804315, 0.04765609, -0.008533617, 0.006844215, -0.026217818, -0.049941003, -0.010676516, -0.039658975,



Tips to scale

Gemini: Video Understanding

Limitations

While Gemini multimodal models are powerful in many multimodal use cases, it's important to understand the limitations of the models:

- **Content moderation:** The models refuse to provide answers on videos that violate our safety policies.
- **Non-speech sound recognition:** The models that support audio might make mistakes recognizing sound that's not speech.
- **High-speed motion:** The models might make mistakes understanding high-speed motion in video due to the fixed 1 frame per second (fps) sampling rate.
- **Transcription punctuation:** *(if using Gemini 1.5 Flash)* The models might return transcriptions that don't include punctuation.

```
#!/bin/bash
```

```
# Create the output folder if it doesn't exist
```

```
mkdir -p small
```

```
files=(*.mp4)
```

```
for i in "${!files[@]}"; do
```

```
file="${files[$i]}"
```

```
filename=$(basename "$file" .mp4)
```

```
output="small/${filename}.mp4"
```

```
gpu=$((i % 2))
```

```
CUDA_VISIBLE_DEVICES="$gpu" ffmpeg -hwaccel cuda -i "$file" -vf
```

```
"fps=1,scale=1024:1024:force_original_aspect_ratio=decrease,pad=1024:1024:(ow-iw)/2:(oh-ih)/2" \
```

```
-c:v h264_nvenc -preset slow -cq 30 -c:a aac -b:a 96k -movflags +faststart "$output" &
```

```
echo "Started processing: $file on GPU $gpu"
```

```
# Wait for processes to complete every two files
```

```
if ((i % 2 == 1)); then
```

```
wait
```

```
fi
```

```
done
```



1. Use FFMPEG-GPU NVIDIA
2. Reduce size and apply FPS 1

**Process 2,6K videos: 31GB to 2.7GB (90% reduction)
with FPS 1. 10x faster speed than CPU.**

Use PostgreSQL and PGVector

Enable the extension (do this once in each database where you want to use it)

```
CREATE EXTENSION vector;
```



Create a vector column with 3 dimensions

```
CREATE TABLE items (id bigserial PRIMARY KEY, embedding vector(3));
```



Insert vectors

```
INSERT INTO items (embedding) VALUES ('[1,2,3]'), ('[4,5,6]');
```



Get the nearest neighbors by L2 distance

```
SELECT * FROM items ORDER BY embedding <-> '[3,1,2]' LIMIT 5;
```



Also supports inner product (<#>), cosine distance (<=>), and L1 distance (<+> , added in 0.7.0)

Note: <#> returns the negative inner product since Postgres only supports Asc order index scans on operators

<https://github.com/pgvector/pgvector>

Thank you!

Linkedin

[linkedin.com/in/yodiaditya](https://www.linkedin.com/in/yodiaditya)

Github

<https://github.com/yodiaditya/shorts>



Google Developer Groups

